

CS 188 Discussion 4:

Games

Kenny Wang (kwkw@berkeley.edu)

Wed Sep 20, 2023

Slides inspired by Sashrika Pandey and Regina Wang

Administrivia

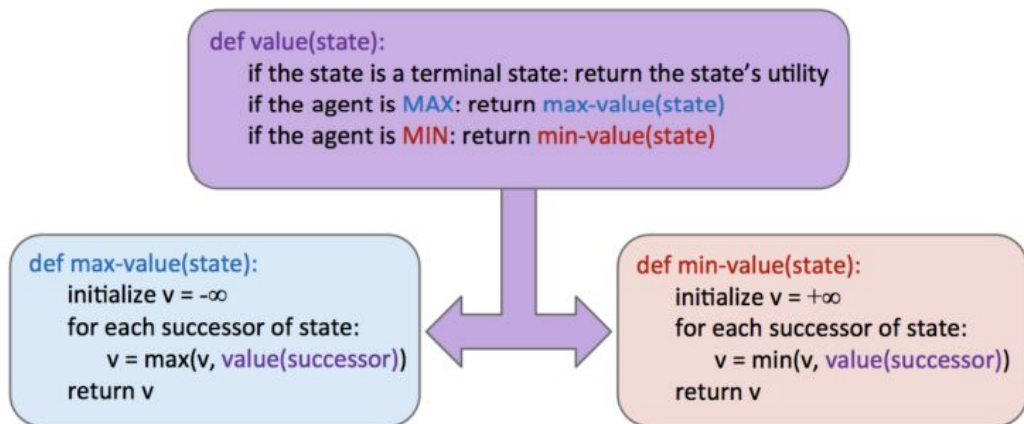
- Project 2 due on Friday, Sep 22
- Homework is due on Tuesdays
- We have office hours pretty much all day every weekday (12-7), come to Soda 341B!
- Reminder: Need extensions? We will give you extensions!

Today's Topics

- Games!
 - Minimax
 - Alpha-beta pruning
 - Expectimax

Minimax


- Zero-sum game where opponent plays optimally
- Agents and opponents take turns
 - Maximizing agents/nodes try to **maximize** utility
 - Minimizer agents/nodes try to **minimize** utility



$$\begin{aligned}\forall \text{ agent-controlled states, } V(s) &= \max_{s' \in \text{successors}(s)} V(s') \\ \forall \text{ opponent-controlled states, } V(s) &= \min_{s' \in \text{successors}(s)} V(s') \\ \forall \text{ terminal states, } V(s) &= \text{known}\end{aligned}$$

Minimax Example

- Pacman maximizes
- Ghost minimizes



Alpha-Beta Pruning

- Minimax runtime is $O(b^m)$:(
 - b = branching factor, m = approx tree depth where terminal nodes are found
- **Alpha-Beta Pruning:**
 - Say we're looking at some node n . When we go through n 's children, if we realize that node n 's value is guaranteed to be made redundant by another node, stop checking n 's children!
 - Runtime about $O(b^{m/2})$

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Alpha-Beta Pruning (continued)

- α represents MAX's best (highest) available option—a minimum value for any MIN nodes
 - if a MIN node has a child $\leq \alpha$, it will never be high enough to change the MAX nodes above!
 - α is updated in MAX nodes to be its current best (highest) value
- β represents MIN's best (lowest) available option—a maximum value for any MAX nodes
 - if a MAX node has a child $\geq \beta$, it will never be small enough to change the MIN nodes above!
 - β is updated in MIN nodes to be its current best (lowest) value
- α and β are passed down to children


α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```


Alpha-Beta Pruning (continued)

- Example for me to copy onto the board



Alpha-Beta Pruning (continued)

- Solution



Worksheet 1(a) + 1(b)


Expectimax

- Minimizer nodes replaced by **chance nodes**, which find the expected value of their children
- Captures non-optimal behavior

$$\forall \text{ agent-controlled states, } V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$\forall \text{ chance states, } V(s) = \sum_{s' \in \text{successors}(s)} p(s'|s) V(s')$$

$$\forall \text{ terminal states, } V(s) = \text{known}$$



Worksheet 1(c) + 1(d)

Worksheet 2

Thank you for attending!

Attendance link:

- <https://tinyurl.com/cs188fa23>

Discussion No: 4

Remember my name is Kenny

My email: kwkw@berkeley.edu

