

CS 188 Discussion 9:

MDP Sampling, Decision Networks, VPI

Kenny Wang (kwkw@berkeley.edu)

Wed Oct 25, 2023

Slides based on Sashrika + Joy

Administrivia

- Project 4 due on Mon, Nov 6
- Homework is due on Tuesdays
- We have office hours pretty much all day every weekday (12-7), come to Soda 341B! (my hours are 1-3 PM on Mondays)
- Discussion slides are on Ed

Today's Topics

- MDP Sampling
 - Prior Sampling
 - Rejection Sampling
 - Likelihood Weighting
 - Gibbs Sampling
- Decision Networks
- Value of Perfect Information (VPI)

Prior Sampling

- Algorithm:

- Generate samples

- For $i=1, 2, \dots, n$ (in topological order)
 - Sample X_i from $P(X_i \mid \text{parents}(X_i))$
- Return (x_1, x_2, \dots, x_n)

- Discard samples inconsistent with evidence
- Calculate probability of the query

- Pros:

- Easy!

- Cons:

- Requires a large number of samples, especially when we condition on unlikely scenarios

Rejection Sampling

Worksheet Q1a

- Algorithm:

- Generate samples

- Input: evidence e_1, \dots, e_k
- For $i=1, 2, \dots, n$
 - Sample x_i from $P(x_i \mid \text{parents}(x_i))$
 - If x_i not consistent with evidence
 - Reject: Return, and no sample is generated in this cycle
- Return (x_1, x_2, \dots, x_n)

- Calculate probability of the query

- Pros:

- Bad samples take less time to create, so more efficient than Prior Sampling

- Cons:

- Still requires a large number of samples, especially when we condition on unlikely scenarios

Likelihood Weighting


Worksheet Q1b


- Idea: Set all your variables to agree with evidence, then sample
 - **Weight** all samples by the probability of the evidence given the sampled variables
- Pros:
 - Never generate a bad sample!
- Cons:
 - Evidence influences the choice of downstream variables, but not upstream ones. Samples could have low weight!

```
▪ Input: evidence  $e_1, \dots, e_k$ 
▪  $w = 1.0$ 
▪ for  $i=1, 2, \dots, n$ 
  ▪ if  $X_i$  is an evidence variable
    ▪  $x_i =$  observed value $i$  for  $X_i$ 
    ▪ Set  $w = w * P(x_i \mid \text{parents}(X_i))$ 
  ▪ else
    ▪ Sample  $x_i$  from  $P(X_i \mid \text{parents}(X_i))$ 
▪ return  $(x_1, x_2, \dots, x_n), w$ 
```

Reinforcement Learning Overview

- We solved MDPs using **offline planning**
 - Knew exact transition and reward functions, could precompute optimal actions without trying anything
- Now we move to **online planning**
 - Agent has no prior knowledge of transitions and rewards
 - Try **exploration** to receive feedback
 - Estimate an optimal policy to use for **exploitation** (maximizing rewards)





Offline Solution:

Compute policy ahead
of time



Online Learning:

Compute policy as
experience comes in

Types of Reinforcement Learning

- **Passive RL:** learn values from experience using a given policy, then use that to inform better policies
 - **Model-based:** learn the MDP model (transitions and rewards) from experiences, then solve the MDP
 - **Model-free:** forego learning the MDP model, directly learn $V(s)$ or $Q(S, a)$
 - **Direct Evaluation**
 - **TD-learning (Temporal Difference)**
- **Active RL:** learn policy from our experiences directly
 - **Q-learning:** learns $Q(\text{state}, \text{action})$ values of the optimal policy
 - **Approximate Q-learning**

Model-Based Learning

- **Model-based learning:** observe a bunch of actions, then estimate transition probabilities $T(s, a, s')$ and rewards $R(s, a, s')$ by averaging.
 - $T(s, a, s')$ is the number of times you end up in s' after being in state s and taking action a , divided by the total number of times you took a from s .
 - $R(s, a, s')$ is the average reward you got from starting in s , taking action a , and ending up in s' .

Worksheet 1(a)

TD-Learning (model-free, passive RL)

- Idea: learn values of states, given that you're following some policy

- **Algorithm**

- Initialize value estimates for policy π $\forall s, V^\pi(s) = 0$

- Each timestep:

- Take action $\pi(s)$ and receive reward $R(s, \pi(s), s')$

- Obtain sample $\text{sample} = R(s, \pi(s), s') + \gamma V^\pi(s')$

- Update value estimate with **exponential moving average**

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha \cdot \text{sample}$$

- Learning rate $\alpha, 0 \leq \alpha \leq 1$

Q-Learning (active RL)

- Idea: learn $Q(s, a)$ values - which gives you the optimal policy (maximize Q)
- **Algorithm**
 - Initialize $Q(\text{state}, \text{action})$ estimates to 0
 - Each timestep:
 - Take some action, any action! It doesn't need to be optimal
 - Obtain sample $\text{sample} = R(s, a, s') + \gamma \max_{a'} Q(s', a')$
 - Update Q-value estimate with exponential moving average
$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \cdot \text{sample}$$
 - If we explore enough and decrease α appropriately, learn optimal Q-values

Worksheet 1(b)

Value Iteration

Value Iteration

- A dynamic programming algorithm we use to compute values until convergence ($\forall s, U_{k+1}(s) = U_k(s)$)
- Algorithm
 - $\forall s \in S$, initialize $U_0(s) = 0$ [initial value estimate is 0]
 - Repeat rule until convergence (U-values stop changing)

$$\forall s \in S, U_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma U_k(s')]$$

- Convergence is when
 - $\forall s \in S, U_k(s) = U_{k+1}(s) = U^*(s)$

Policy Iteration

Policy Iteration

- Issues with value iteration
 - $O(|S|^2|A|)$ runtime
 - Overcomputes, policy tends to converge faster than values
- Policy iteration: preserve the optimality from value iteration but with better performance by iterating until only the *policy* converges instead of the U-values

Policy Iteration

- Algorithm

- Define initial policy π_0 (can be arbitrary)
- Repeat until convergence:
 - **Policy evaluation:** compute expected utility of starting in state s when following policy π , for all states s

$$U^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma U^\pi(s')]$$

- **Policy improvement:** generate a better policy

$$\pi_{i+1}(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma U^{\pi_i}(s')]$$

- Convergence when $\pi_{i+1} = \pi_i$ (policy stops changing)

Summary

- **Value Iteration**

- $\forall s \in S$, initialize $U_0(s) = 0$
[initial value estimate is 0]
- Repeat rule until convergence
(U-values stop changing)

$$\forall s \in S, U_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma U_k(s')]$$

- Convergence is when
 $\forall s \in S, U_k(s) = U_{k+1}(s) = U^*(s)$

- **$Q^*(s,a)$: the optimal value of (s, a)** [state, action pair]

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma U^*(s')]$$

- **$U^*(s)$: the optimal value of state s** $U^*(s) = \max_a Q^*(s, a)$

- **Policy Iteration**

- Define initial policy π_0 (can be arbitrary)
- Repeat until convergence:
 - **Policy evaluation:** compute expected utility of starting in state s when following policy π , for all states s

$$U^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma U^\pi(s')]$$

- **Policy improvement:** generate a better policy

$$\pi_{i+1}(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma U^{\pi_i}(s')]$$

- Convergence when $\pi_{i+1} = \pi_i$ (policy stops changing)

Rest of the Worksheet

Thank you for attending!

Attendance link:

- <https://tinyurl.com/cs188fa23>

Discussion No: 6

Remember my name is Kenny

My email: kwkw@berkeley.edu

