

Last discussion! (except for final review discussion)

CS 188 Discussion 12:

Logistic Regression, Neural Networks

Kenny Wang (kwkw@berkeley.edu)

Wed Nov 15, 2023

Slides based on Sashrika + Joy

Administrivia

- Project 4 deadline extended Mon, Nov 6 => Fri, Nov 10
- Homework is due on Tuesdays
- We have office hours pretty much all day every weekday (12-7), come to Soda 341B! (my hours are 1-3 PM on Mondays)
- Discussion slides are on Ed

Today's Topics

- Sigmoid Function
- Logistic Regression
- Gradient Descent
- Neural Networks!!


Sigmoid Function

mildly important in CS 189!

- **Sigmoid (aka Logistic) function:**

$$\phi(z) = \frac{1}{1 + e^{-z}}$$

- All outputs bounded by 0 and 1, is 0.5 when z is 0
- Can convert linear outputs to probabilities (logistic regression)
- Used to introduce nonlinearity in neural networks




Logistic Regression

- **Idea:** Instead of simply using $w^T x$, apply sigmoid function on $w^T x$

- **Logistic Regression**

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

- Can predict probabilities, unlike binary perceptron



- How to compute \mathbf{w} ? Use gradient descent!
- See also: multiclass logistic regression, using the **softmax function**

Optimization: Gradient Ascent/Descent

- **Goal:** Want to find the parameters that maximize **objective function** or minimize **loss function**
- If closed-form formula for global optimum does not exist, can use **gradient ascent/descent**
- Observation: gradient is direction of steepest increase ... by repeatedly following the gradient, we can chase maxima/minima

- **Gradient Ascent**

```
Randomly initialize  $\mathbf{w}$   
while  $\mathbf{w}$  not converged do  
  |  $\mathbf{w} \leftarrow \mathbf{w} + \alpha \nabla_{\mathbf{w}} f(\mathbf{w})$   
end
```

- **Gradient Descent**

```
Randomly initialize  $\mathbf{w}$   
while  $\mathbf{w}$  not converged do  
  |  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla_{\mathbf{w}} f(\mathbf{w})$   
end
```

Learning rate α is a
hyperparameter

Musheen Lurning!!1


- **Core Idea:** We give machines access to data and they learn for themselves!
- Data is often split into training, validation, and test sets
 - **Training set:** Used to fit the model
 - **Validation set:** Used to tune **hyperparameters** (learning rate, model structure, etc)
 - **Test set:** Used to test the entire model
- Some types of machine learning problems
 - **Classification problems:** try to classify data into discrete classes
 - **Regression problems:** try to estimate some numerical value from data
 - **Clustering problems:** try to group similar data into clusters
- Types of learning
 - **Supervised learning:** training data has labels, e.g. classification
 - **Unsupervised learning:** training data has no labels, e.g. clustering

Naive Bayes

- **Goal:** create a model that can predict a label Y given features, where we assume all features are independently affected by the label
- Ex: Spam filter
 - Y is in {Spam, Ham}
 - F_i in {0, 1} is whether word i appears in the email.
- Label email based on the higher of these two probabilities:

$$P(Y = ham|F_1 = f_1, \dots, F_n = f_n) \quad P(Y = spam|F_1 = f_1, \dots, F_n = f_n)$$

- Generalized:
$$\begin{aligned} prediction(f_1, \dots, f_n) &= \underset{y}{\operatorname{argmax}} P(Y = y|F_1 = f_1, \dots, F_n = f_n) \\ &= \underset{y}{\operatorname{argmax}} P(Y = y, F_1 = f_1, \dots, F_n = f_n) \\ &= \underset{y}{\operatorname{argmax}} P(Y = y) \prod_{i=1}^n P(F_i = f_i|Y = y) \end{aligned}$$




Maximum Likelihood Estimation

- How to estimate CPTs, since we don't actually know them?
 - Parameter estimation with MLE
- Find the probabilities (CPT values) $\theta = \mathbf{P}(\cdot)$ such that we maximize the likelihood of observing our observations, $P(\text{observations} \mid \theta)$
- Answer is actually fairly intuitive. f/e, given data (F, Y) ,

$$\begin{aligned} P(Y=y) &= \text{MLE}(\theta \mid (F, Y)) \\ &= \frac{\text{\# examples with } Y=y}{\text{total \# examples}} \end{aligned}$$

$$\begin{aligned} P(F_i=f \mid Y=y) &= \text{MLE}(\theta \mid (F, Y)) \\ &= \frac{\text{\# examples with } (F_i=f, Y=y)}{\text{\# examples with } Y=y} \end{aligned}$$



Laplace Smoothing

Worksheet Q1

- Chance of **overfitting** (model doesn't generalize well post-training) with our parameter estimation
- **Laplace Smoothing**: pretend you saw each of the $|X|$ possible outcomes k extra times

$$P_{LAP,k}(x) = \frac{\text{count}(x) + k}{N + k|X|}$$

$$P_{LAP,k}(x|y) = \frac{\text{count}(x,y) + k}{\text{count}(y) + k|X|}$$


- k is a **hyperparameter**, meaning you can choose what to set it to
 - Smaller k means your probability estimates follow the training data more closely
 - Larger k means your probability estimates are more uniform

Perceptrons

- Naive Bayes
- **Perceptrons**

Binary Perceptrons

- Idea: Linearly separate data into two classes using a decision boundary (defined by a set of weights)
- If the data is linearly separable, the algorithm will perfectly classify the data!
- To find boundaries that don't have to cross the origin, incorporate a “bias” feature that always has value 1



Perceptron Algorithm

Worksheet Q2

1. Initialize all weights to 0: $\mathbf{w} = \mathbf{0}$
2. For each training sample, with features $\mathbf{f}(x)$ and true class label $y^* \in \{-1, +1\}$, do:
 - (a) Classify the sample using the current weights, let y be the class predicted by your current \mathbf{w} :

$$y = \text{classify}(x) = \begin{cases} +1 & \text{if } \text{activation}_w(x) = \mathbf{w}^T \mathbf{f}(x) > 0 \\ -1 & \text{if } \text{activation}_w(x) = \mathbf{w}^T \mathbf{f}(x) < 0 \end{cases}$$

- (b) Compare the predicted label y to the true label y^* :
 - If $y = y^*$, do nothing
 - Otherwise, if $y \neq y^*$, then update your weights: $\mathbf{w} \leftarrow \mathbf{w} + y^* \mathbf{f}(x)$
3. If you went through **every** training sample without having to update your weights (all samples predicted correctly), then terminate. Else, repeat step 2.

Rest of the Worksheet

Thank you for attending!

Attendance link:

- <https://tinyurl.com/cs188fa23>

Discussion No: 11

Remember my name is Kenny

My email: kwkw@berkeley.edu

